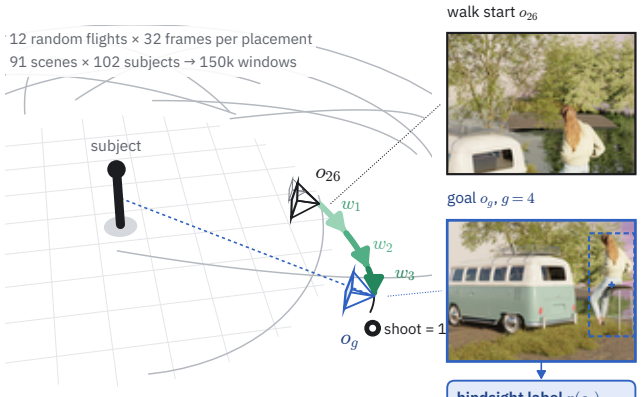
(a) Self-labelling data engine

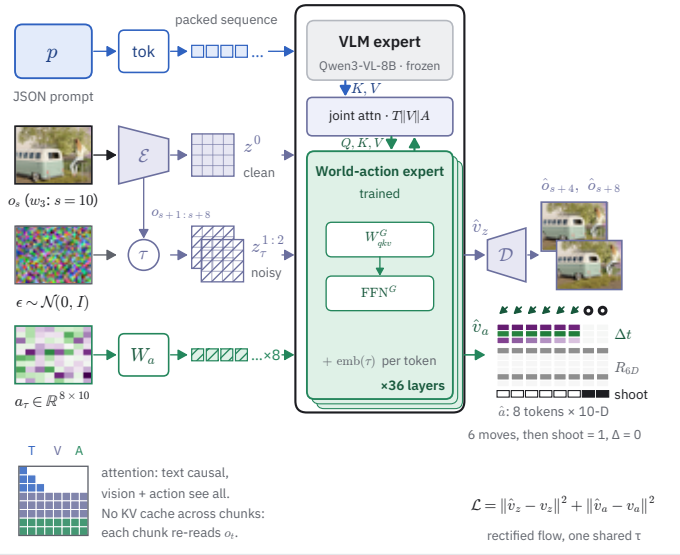


one goal → three 8-step windows w_1, w_2, w_3

$$(o_s,\ p(o_g)) \mapsto (a_{s:s+8},\ o_{s+1:s+8}), \quad |g-s| \le 32$$

from the projected mesh: no reward, no human label

(b) World action model



(c) Deployment: intent → shot profile → closed-loop motion

the goal image is never shown to the policy

